

# Turning the PageRank on PDN

Lexie Barthelemess  
CS Senior

Mark Cunningham  
CS Senior

Quentin Windsor  
CS Senior

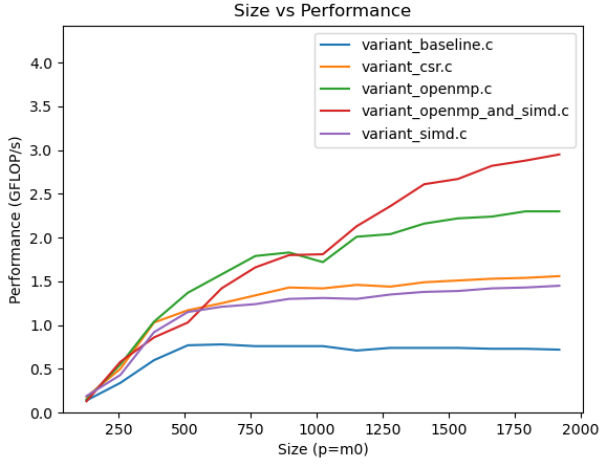


Fig. 1. All tested variants

**Abstract**—The PageRank algorithm was developed to rank web pages for search engines. A page is considered “important” if other important pages link to it. Each page is represented as a node on a graph. This paper examines storing these nodes in different sparse data formats as well as different parallel constructs for iteration to convergence.

## I. COO TO CSR

The CSR memory format is more space-efficient than COO since we are able to consolidate references to the row index. Thus, we need to store less values for row in comparison to COO, which stores values 1-for-1-for-1.

In the transformation function that we wrote to take a COO matrix and produce a CSR, we first allocate appropriate memory for the CSR matrix and copy the metadata over. Then, it allocates buffers for row, column, and values. These are filled with information from the COO matrix.

The CSR format is beneficial for performance since it optimizes memory usage through contiguous memory access; since the non-zero elements are stored consecutively, we can avoid cache misses and improve efficiency for operations such as this matrix-vector multiplication.

Figure 1 shows the difference between using the baseline (COO) format and using the CSR format.

## II. SIMD

To leverage repeated instances of consecutive row indices, we concurrently processed eight rows at a time where possible. To eliminate bounds checking during matrix multiplication, we divided the row index array into two groups: SIMD and regular. Each group is an array containing (start, end) indices. In the

SIMD version, pairs spanned eight elements, while the regular version encompassed all other pairs. As the number of non zero entries grew, the percentage of non-zero entries covered by the SIMD operation increased. About 68% of entries are performed by SIMD at size 128, and about 89% at size 1024. The percent performed by SIMD levels off at around 92%.

The matrix vector multiplication was split into two for-loops: One that spans the size of the SIMD entries and the other spanning the size of the regular entries. In each iteration of the SIMD loop, eight column indices are loaded from the column-index-array, which are used to gather the corresponding values from the x-array. Then, eight values are loaded from the values-array and are multiplied by the values from the x-array. The resulting vector is summed into a single float value which is added to the value at the current index of the y-array.

The SIMD variant nearly doubled the performance across all sizes when compared to the baseline, but did not see further improvements as the size increased.

## III. OPENMP

To parallelize the matrix vector multiplication, we used openMP to split the work among threads in a parallel region. Since there were many occasions where it was possible for two threads to be writing to the same element in the y-array concurrently, we made sure that each thread had its own private copy of the y-array. We used the reduction clause to avoid race conditions without sacrificing too much performance. After the parallel region is completed, the private copies are copied back to the y-array. The openMP variant performed better than SIMD, and did not completely fall off as size increased, suggesting that it could continue increasing with size.

## IV. SIMD AND OPENMP COMBINED

We chose to combine SIMD and openMP for our 2 Form variant. The performance gain was roughly the same as openMP by itself for sizes less than 1000, but began to show a noticeable improvement past sizes greater than 1000.

## V. CONCLUSION

All variants showed improvement when compared to the baseline. Data access patterns are crucial to implementing successful parallel code, and the PageRank algorithm is no exception. Optimizing data access patterns in the context of PageRank involved minimizing data dependencies, maximizing locality of data, and ensuring balanced workloads among parallel constructs. If we had more time, we would have implemented CSR in combination with the other variants.